

Porter Case Study

Introduction

- Porter, India's largest marketplace for intra-city logistics, is revolutionizing the delivery sector with technology-driven solutions.
- This case focuses on leveraging neural networks to accurately predict delivery times, a critical aspect of customer satisfaction in logistics.
- With a dataset encompassing various aspects of orders and deliveries, Porter aims to refine its delivery time estimations.
- Analyzing this dataset can provide significant insights into delivery dynamics, efficiency bottlenecks, and optimization opportunities.
- The insights obtained can enhance Porter's operational efficiency, ensuring timely deliveries and improving driver-partner allocation.

What is expected?

- As a data scientist at Porter, your task is to analyze the dataset to accurately predict delivery times for different orders. Your primary goal is to build a regression model using neural networks, evaluate its performance, and provide insights for optimizing delivery operations.

1. Data

The analysis was done on the data located at -

https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/015/039/original/dataset.csv.zip
1663710760

2. Libraries

Below are the libraries required

In [105...

```
# Libraries to analyze data
import numpy as np
import pandas as pd

# Libraries to visualize data
import matplotlib.pyplot as plt
```

```

import seaborn as sns

from sklearn.neighbors import LocalOutlierFactor

from sklearn.impute import KNNImputer

from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation, LeakyReLU, Dropout, BatchNormalization
from tensorflow.keras.optimizers import Adam, RMSprop
from sklearn.metrics import mean_squared_error

```

3. Data Loading

Loading the data into Pandas dataframe for easily handling of data

```

In [2]: # read the file into a pandas dataframe
df = pd.read_csv('dataset.csv')
# Look at the datatypes of the columns
print('*****')
print(df.info())
print('*****\n')
print('*****')
print(f'Shape of the dataset is {df.shape}')
print('*****\n')
print('*****')
print(f'Number of nan/null values in each column: \n{df.isna().sum()}')
print('*****\n')
print('*****')
print(f'Number of unique values in each column: \n{df.nunique()}')
print('*****\n')
print('*****')
print(f'Duplicate entries: \n{df.duplicated().value_counts()}')

```

```

*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            196441 non-null  float64
1   created_at                           197428 non-null  object
2   actual_delivery_time                 197421 non-null  object
3   store_id                             197428 non-null  object
4   store_primary_category               192668 non-null  object
5   order_protocol                       196433 non-null  float64
6   total_items                          197428 non-null  int64
7   subtotal                             197428 non-null  int64
8   num_distinct_items                  197428 non-null  int64
9   min_item_price                      197428 non-null  int64
10  max_item_price                       197428 non-null  int64
11  total_onshift_partners               181166 non-null  float64
12  total_busy_partners                  181166 non-null  float64
13  total_outstanding_orders             181166 non-null  float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB
None

```

```

*****
Shape of the dataset is (197428, 14)
*****

```

```

*****
Number of nan/null values in each column:
market_id                987
created_at                0
actual_delivery_time      7
store_id                  0
store_primary_category    4760
order_protocol            995
total_items               0
subtotal                  0
num_distinct_items        0
min_item_price            0
max_item_price            0
total_onshift_partners    16262
total_busy_partners       16262
total_outstanding_orders  16262
dtype: int64
*****

```

```

*****
Number of unique values in each column:
market_id                6
created_at               180985
actual_delivery_time     178110
store_id                  6743
store_primary_category    74
order_protocol            7

```

```

total_items          57
subtotal            8368
num_distinct_items   20
min_item_price       2312
max_item_price       2652
total_onshift_partners 172
total_busy_partners   159
total_outstanding_orders 281
dtype: int64
*****

*****

Duplicate entries:
False    197428
Name: count, dtype: int64

```

```

In [3]: # Look at the top 20 rows
df.head(5)

```

```

Out[3]:

```

	market_id	created_at	actual_delivery_time	store_id	store_pi
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	

```

In [4]: df.describe()

```

Out[4]:	market_id	order_protocol	total_items	subtotal	num_distinct_items	m
count	196441.000000	196433.000000	197428.000000	197428.000000	197428.000000	1
mean	2.978706	2.882352	3.196391	2682.331402	2.670791	
std	1.524867	1.503771	2.666546	1823.093688	1.630255	
min	1.000000	1.000000	1.000000	0.000000	1.000000	
25%	2.000000	1.000000	2.000000	1400.000000	1.000000	
50%	3.000000	3.000000	3.000000	2200.000000	2.000000	
75%	4.000000	4.000000	4.000000	3395.000000	3.000000	
max	6.000000	7.000000	411.000000	27100.000000	20.000000	

In [5]: `df.describe(include='object')`

Out[5]:	created_at	actual_delivery_time	store_id	store_primary_
count	197428	197421	197428	
unique	180985	178110	6743	
top	2015-02-11 19:50:43	2015-02-11 20:40:45	d43ab110ab2489d6b9b2caa394bf920f	
freq	6	5	937	

In [6]: `min(df['created_at']), max(df['created_at'])`

Out[6]: ('2014-10-19 05:24:15', '2015-02-18 06:00:44')

Insight

- There are **197428** entries with 14 columns
- The data is available between **19-Oct-2014 to 18-Feb-2015**, around 5 months of data
- There are null/missing values in each of the dates
- There are no **duplicates**
- The columns **market_id**, **total_onshift_partners**, **total_busy_partners** and **total_outstanding_orders** can be of type **int64**
- The columns **market_id**, **order_protocol** and **num_distinct_items** can be converted to categorical columns
- The columns **created_at** and **actual_delivery_time** need to be of type **datetime**
- Extract **hour** and **day** of the order placement from **created_at**
- Create **delivery_time_mins** column by subtracting **created_at** from **actual_delivery_time**
- Columns **created_at** and **actual_delivery_time** can be dropped

```
In [7]: df[['total_onshift_partners', 'total_busy_partners', 'total_outstanding_orders']]
df[['market_id', 'order_protocol', 'num_distinct_items']] = df[['market_id', 'order
df[['created_at', 'actual_delivery_time']] = df[['created_at', 'actual_delivery_tim
df['created_hour']=df['created_at'].dt.hour
df['created_day']=df['created_at'].dt.dayofweek
df['delivery_time_mins'] = round((df['actual_delivery_time'] - df['created_at'])/pd
df.drop(columns=['created_at', 'actual_delivery_time'], inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 197428 entries, 0 to 197427
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	market_id	196441 non-null	category
1	store_id	197428 non-null	object
2	store_primary_category	192668 non-null	object
3	order_protocol	196433 non-null	category
4	total_items	197428 non-null	int64
5	subtotal	197428 non-null	int64
6	num_distinct_items	197428 non-null	category
7	min_item_price	197428 non-null	int64
8	max_item_price	197428 non-null	int64
9	total_onshift_partners	181166 non-null	Int64
10	total_busy_partners	181166 non-null	Int64
11	total_outstanding_orders	181166 non-null	Int64
12	created_hour	197428 non-null	int32
13	created_day	197428 non-null	int32
14	delivery_time_mins	197421 non-null	float64

```
dtypes: Int64(3), category(3), float64(1), int32(2), int64(4), object(2)
```

```
memory usage: 17.7+ MB
```

```
In [8]: df.describe()
```

Out[8]:

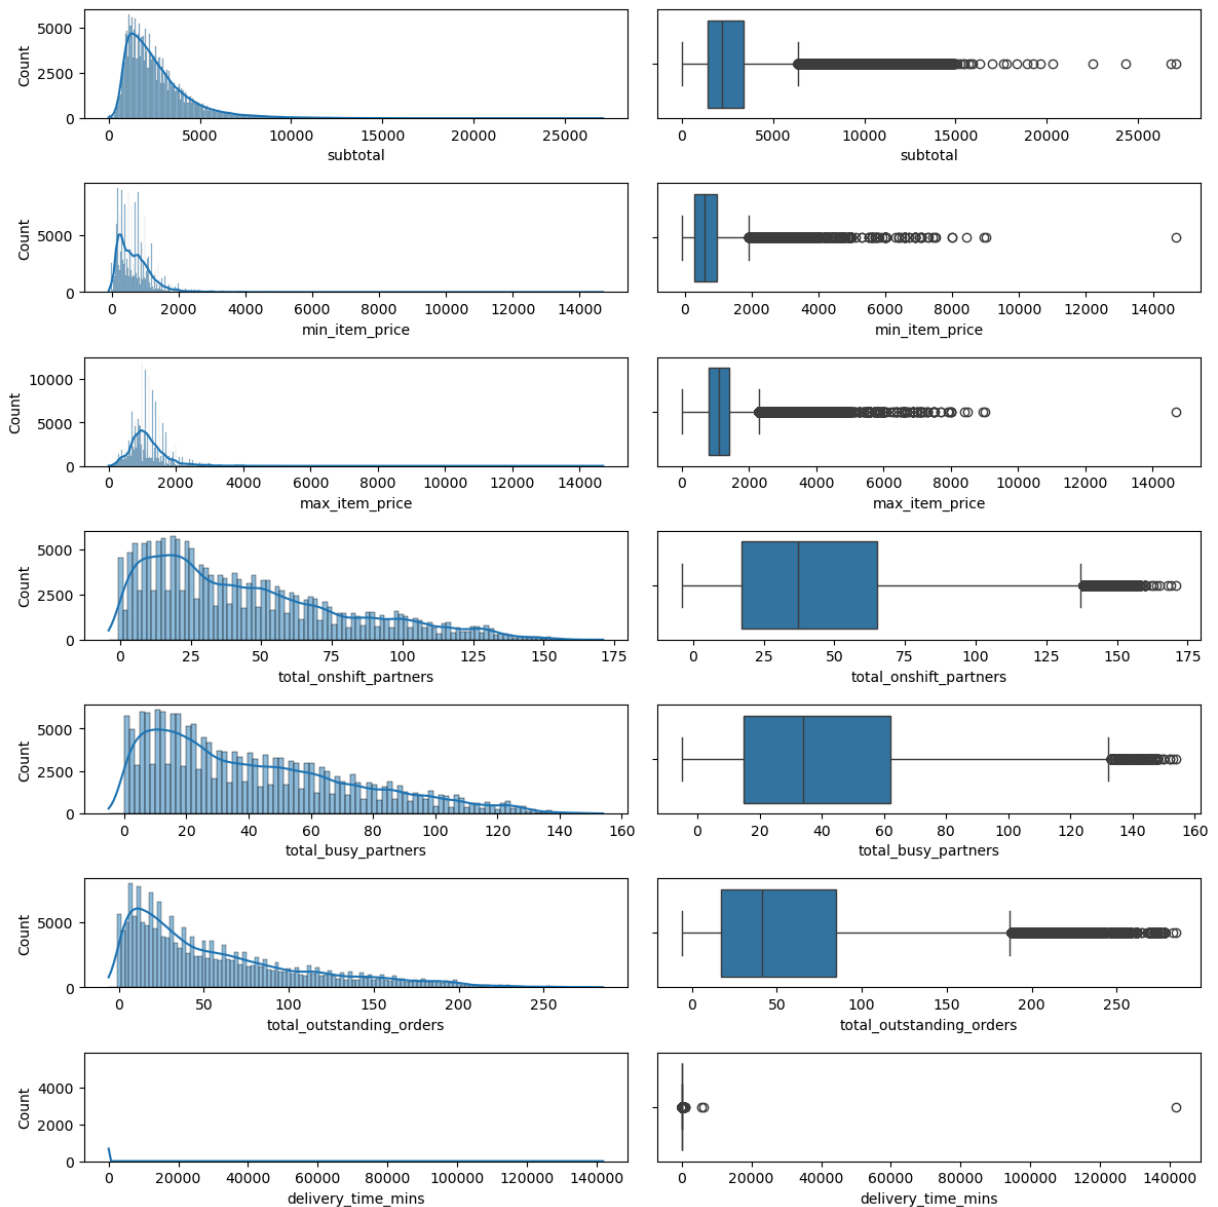
	total_items	subtotal	min_item_price	max_item_price	total_onshift_partners
count	197428.000000	197428.000000	197428.000000	197428.000000	181166.0
mean	3.196391	2682.331402	686.218470	1159.588630	44.808093
std	2.666546	1823.093688	522.038648	558.411377	34.526783
min	1.000000	0.000000	-86.000000	0.000000	-4.0
25%	2.000000	1400.000000	299.000000	800.000000	17.0
50%	3.000000	2200.000000	595.000000	1095.000000	37.0
75%	4.000000	3395.000000	949.000000	1395.000000	65.0
max	411.000000	27100.000000	14700.000000	14700.000000	171.0

4. Exploratory Data Analysis

4.1. Univariate Analysis

4.1.1. Numerical variables

```
In [9]: fig, axes = plt.subplots(nrows=7, ncols=2, figsize = (12, 12))
sns.histplot(data=df, x = "subtotal", kde=True, ax=axes[0,0])
sns.boxplot(data=df, x = "subtotal", ax=axes[0,1])
sns.histplot(data=df, x = "min_item_price", kde=True, ax=axes[1,0])
sns.boxplot(data=df, x = "min_item_price", ax=axes[1,1])
sns.histplot(data=df, x = "max_item_price", kde=True, ax=axes[2,0])
sns.boxplot(data=df, x = "max_item_price", ax=axes[2,1])
sns.histplot(data=df, x = "total_onshift_partners", kde=True, ax=axes[3,0])
sns.boxplot(data=df, x = "total_onshift_partners", ax=axes[3,1])
sns.histplot(data=df, x = "total_busy_partners", kde=True, ax=axes[4,0])
sns.boxplot(data=df, x = "total_busy_partners", ax=axes[4,1])
sns.histplot(data=df, x = "total_outstanding_orders", kde=True, ax=axes[5,0])
sns.boxplot(data=df, x = "total_outstanding_orders", ax=axes[5,1])
sns.histplot(data=df, x = "delivery_time_mins", kde=True, ax=axes[6,0])
sns.boxplot(data=df, x = "delivery_time_mins", ax=axes[6,1])
plt.tight_layout()
plt.show()
```



Insight

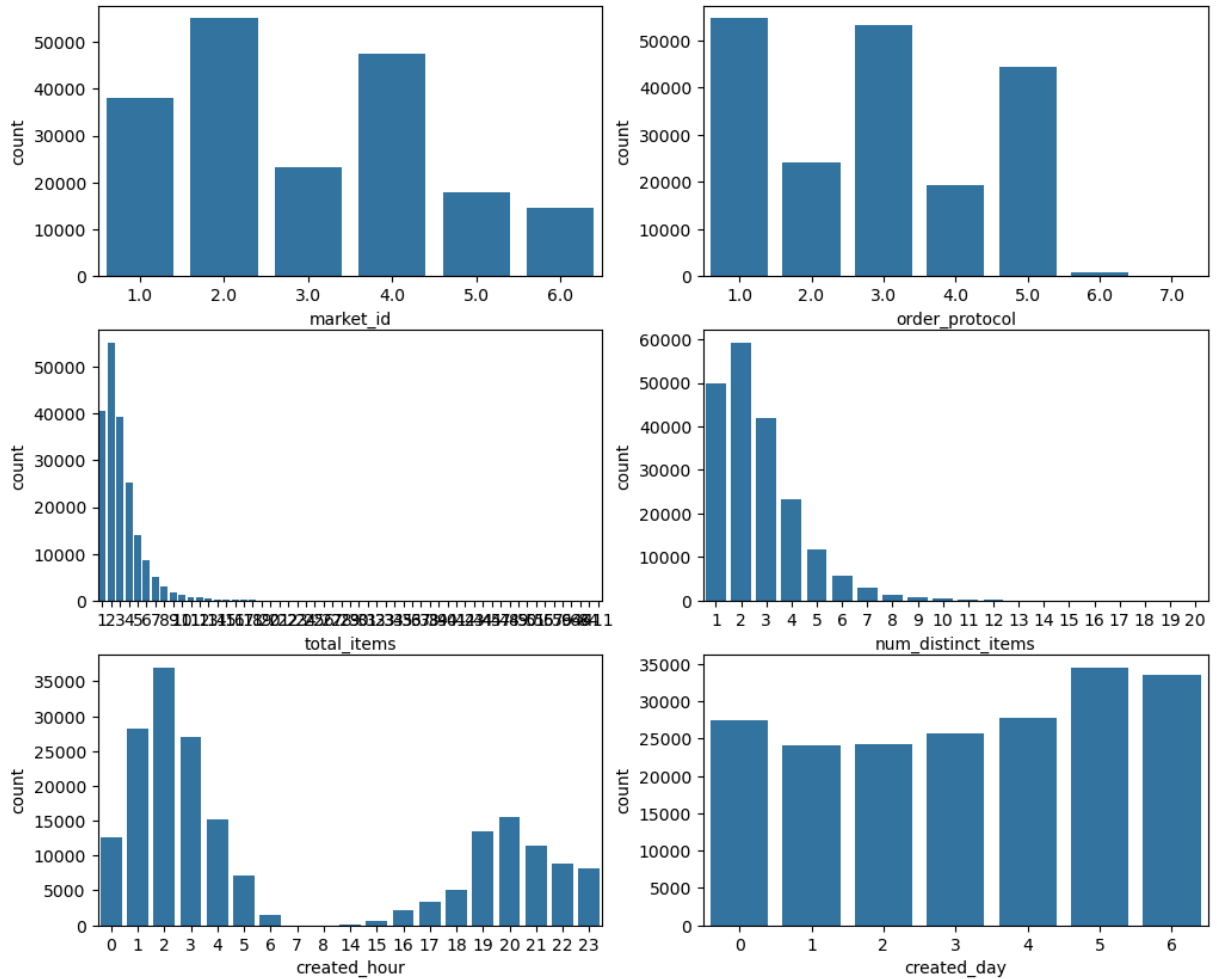
- Majority of the **subtotal** is in the range of 0 to 5000
- Majority of **min_item_price** and **max_item_price** are in the range of 0 to 2000
- **total_onshift_partners**, **total_busy_partners**, **total_outstanding_orders** and **delivery_time_mins** seem to follow similar kind of distribution - right skewed
- The boxplot clearly shows the presence of **outliers** in **subtotal**, **min_item_price**, **max_item_price** and **delivery_time_mins**

4.1.2. Categorical variables

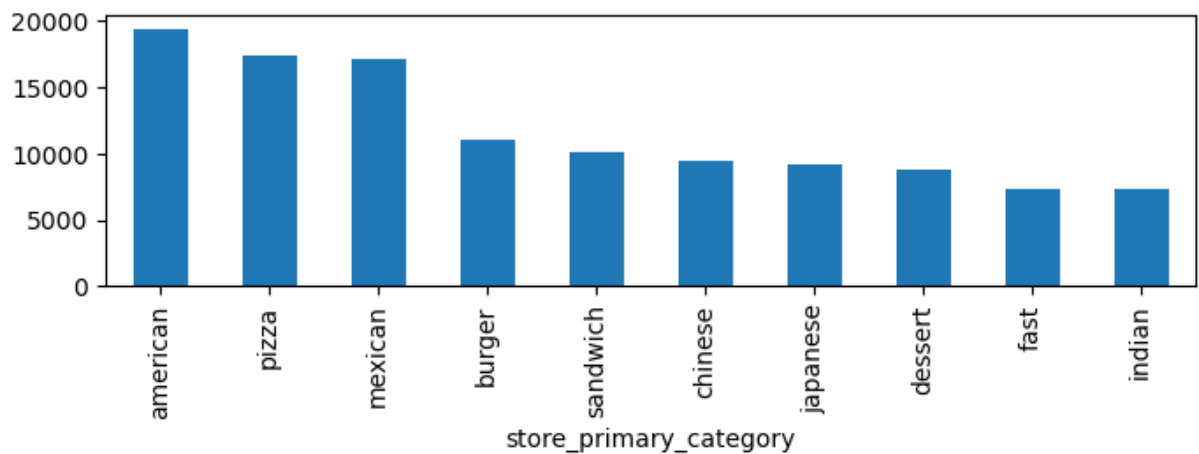
```
In [10]: fig, axes = plt.subplots(3,2,figsize=(12,10))
sns.countplot(ax=axes[0,0], data=df, x='market_id')
sns.countplot(ax=axes[0,1], data=df, x='order_protocol')
sns.countplot(ax=axes[1,0], data=df, x='total_items')
```



```
sns.countplot(ax=axes[1,1], data=df, x='num_distinct_items')
sns.countplot(ax=axes[2,0], data=df, x='created_hour')
sns.countplot(ax=axes[2,1], data=df, x='created_day')
plt.show()
```



```
In [11]: df['store_primary_category'].value_counts()[:10].plot(kind='bar', figsize = (8, 2))
plt.show()
```



```
In [12]: df['store_id'].value_counts()[:5]
```

```
Out[12]: store_id
d43ab110ab2489d6b9b2caa394bf920f    937
757b505cfd34c64c85ca5b5690ee5293    863
faacbcd5bf1d018912c116bf2783e9a1    815
cfecdb276f634854f3ef915e2e980c31    765
45c48cce2e2d7fbdea1afc51c7c6ad26    721
Name: count, dtype: int64
```

Insight

- market_id **2** is the major contributor
- Majority of order placement are through protocol **1** and **3**
- Majority of orders have **2** num_distinct_items as well as **2** total_items
- Majority of the orders are placed at around **2AM**
- Majority of the orders are placed on **weekends**
- Majority of orders delivered are from **american** restaurant
- Majority of orders delivered are from **d43ab110ab2489d6b9b2caa394bf920f** store

4.2. Missing value treatment

```
In [13]: df.isna().sum()/len(df)*100
```

```
Out[13]: market_id          0.499929
store_id          0.000000
store_primary_category  2.411006
order_protocol     0.503981
total_items        0.000000
subtotal           0.000000
num_distinct_items  0.000000
min_item_price      0.000000
max_item_price      0.000000
total_onshift_partners  8.236927
total_busy_partners   8.236927
total_outstanding_orders 8.236927
created_hour         0.000000
created_day          0.000000
delivery_time_mins    0.003546
dtype: float64
```

Insight

- Only **0.5%** of data has missing market_id. I will **drop** all these entries
- **2.4%** of data has missing store_primary_category. I will **replace** these with 'other' category
- Only **0.5%** of data has missing order_protocol. I will **drop** all these entries
- **8.24%** of data has missing total_onshift_partners, total_busy_partners and total_outstanding_orders each. I will use **KNN imputation** to replace the missing value.
- Only **0.0035%** of data has missing delivery_time_mins. I will **drop** these entries too

```
In [14]: df.dropna(subset=['market_id', 'order_protocol', 'delivery_time_mins'], inplace=True)
df.fillna({'store_primary_category':'other'}, inplace=True)
```

Remove rows which have -ve values in min_item_price, max_item_price, total_onshift_partners, total_busy_partners, total_outstanding_orders

```
In [15]: mask = (df['min_item_price'] >= 0) & (df['max_item_price'] >= 0) & (df['total_onshi
df = df[mask]
```


```
In [16]: # Columns to apply KNN Imputation
columns_to_impute = ['total_onshift_partners', 'total_busy_partners', 'total_outsta
# Initialize KNNImputer
imputer = KNNImputer(n_neighbors=2)
# Apply KNN Imputer only to specified columns
df[columns_to_impute] = imputer.fit_transform(df[columns_to_impute])
```

```
In [17]: df.isna().sum()/len(df)*100
```

```
Out[17]: market_id          0.0
store_id          0.0
store_primary_category  0.0
order_protocol    0.0
total_items       0.0
subtotal          0.0
num_distinct_items 0.0
min_item_price    0.0
max_item_price    0.0
total_onshift_partners 0.0
total_busy_partners 0.0
total_outstanding_orders 0.0
created_hour      0.0
created_day       0.0
delivery_time_mins 0.0
dtype: float64
```

Let us look at the distribution again after imputation

```
In [18]: fig, axes = plt.subplots(nrows=3, ncols=2, figsize = (12, 6))
sns.histplot(data=df, x = "total_onshift_partners", kde=True, ax=axes[0,0])
sns.boxplot(data=df, x = "total_onshift_partners", ax=axes[0,1])
sns.histplot(data=df, x = "total_busy_partners", kde=True, ax=axes[1,0])
sns.boxplot(data=df, x = "total_busy_partners", ax=axes[1,1])
sns.histplot(data=df, x = "total_outstanding_orders", kde=True, ax=axes[2,0])
sns.boxplot(data=df, x = "total_outstanding_orders", ax=axes[2,1])
plt.tight_layout()
plt.show()
```

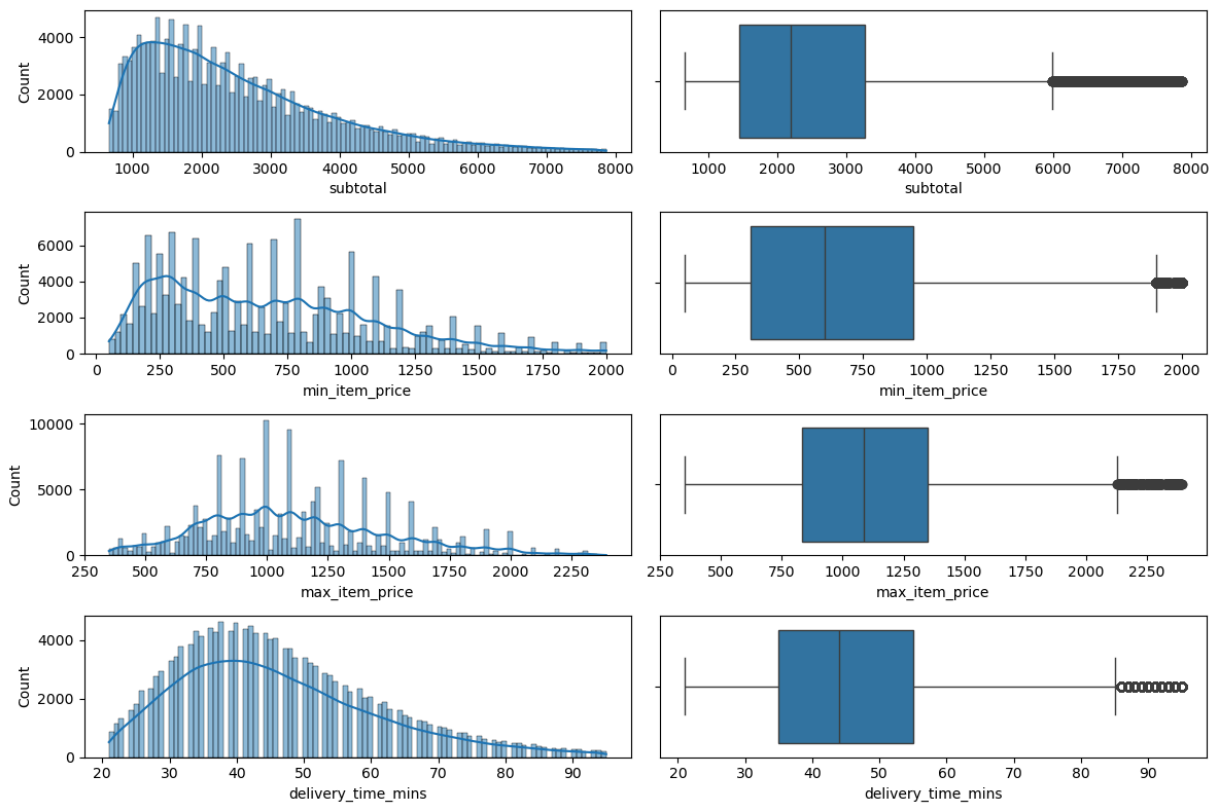


4.2. Outliers treatment


```
In [19]: df['store_primary_category']=df['store_primary_category'].astype('category').cat.codes
df['store_id']=df['store_id'].astype('category').cat.codes
```

```
In [26]: for col in ['subtotal', 'min_item_price', 'max_item_price', 'delivery_time_mins', '
lower_bound = df[col].quantile(0.01)
upper_bound = df[col].quantile(0.99)
df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]
```

```
In [29]: fig, axes = plt.subplots(nrows=4, ncols=2, figsize = (12, 8))
sns.histplot(data=df, x = "subtotal", kde=True, ax=axes[0,0])
sns.boxplot(data=df, x = "subtotal", ax=axes[0,1])
sns.histplot(data=df, x = "min_item_price", kde=True, ax=axes[1,0])
sns.boxplot(data=df, x = "min_item_price", ax=axes[1,1])
sns.histplot(data=df, x = "max_item_price", kde=True, ax=axes[2,0])
sns.boxplot(data=df, x = "max_item_price", ax=axes[2,1])
sns.histplot(data=df, x = "delivery_time_mins", kde=True, ax=axes[3,0])
sns.boxplot(data=df, x = "delivery_time_mins", ax=axes[3,1])
plt.tight_layout()
plt.show()
```



```
In [32]: sns.countplot(data=df, x='total_items')
plt.show()
```




Insight

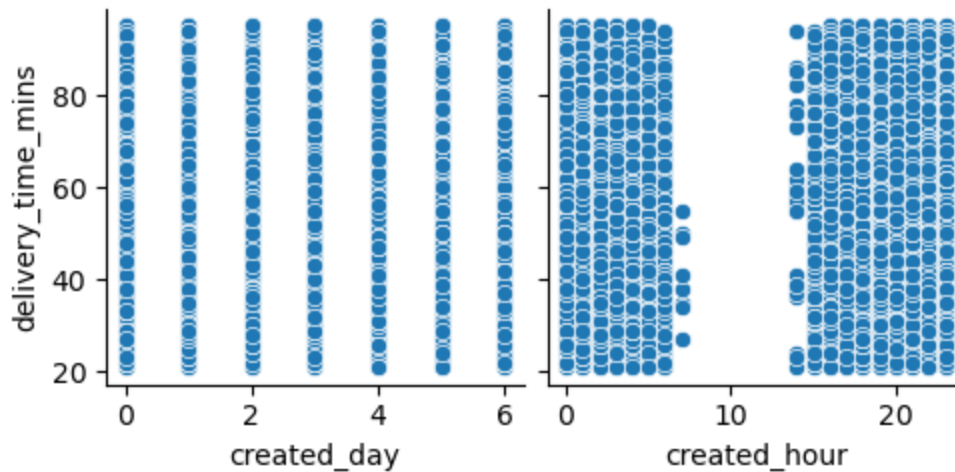
- The distribution of data looks better after removal of outliers
- It can be seen that most of the delivery is done in **40mins**

4.3. Bivariate Analysis

```
In [33]: sns.pairplot(
    df,
    x_vars=['total_items', 'subtotal', 'min_item_price', 'max_item_price', 'total_o
    y_vars=['delivery_time_mins'],)
plt.show()
```




```
In [34]: sns.pairplot(
    df,
    x_vars=['created_day', 'created_hour'],
    y_vars=['delivery_time_mins'],)
plt.show()
```




Insight

- There doesn't seem to be any relationship of delivery time with other features
- Delivery time is almost the same for all days

```
In [35]: sns.heatmap(df.corr())
plt.show()
```



```
In [36]: sns.heatmap(df[['market_id', 'store_id', 'store_primary_category', 'order_protocol',
plt.show()
```



Insight

- Here also we see that there is no relation between delivery time and other features

5. NN Modelling

5.1. Training a NN model

I feel the store id, store primary category and order protocol will not effect delivery time, hece dropping them from feature list

```
In [37]: df_reduced = df.drop(['store_id', 'store_primary_category', 'order_protocol'], axis=1)
```

```
In [39]: # Example to demonstrate categorical encoding
def encode_categorical_data(dataframe):
    # Assume some features are categorical
    categorical_columns = ['market_id']

    # Apply One-Hot Encoding
```



```

one_hot_encoder = OneHotEncoder(sparse_output=False, drop='first') # Drop first
encoded_features = pd.DataFrame(
    one_hot_encoder.fit_transform(dataframe[categorical_columns]),
    columns=one_hot_encoder.get_feature_names_out(categorical_columns)
)

# Drop original categorical columns and concatenate encoded columns
dataframe = dataframe.drop(categorical_columns, axis=1)
dataframe = pd.concat([dataframe.reset_index(drop=True), encoded_features.reset_index(drop=True)], axis=1)
return dataframe

```

```

In [40]: # Encode categorical data
df_reduced_encoded = encode_categorical_data(df_reduced)

```

```

In [41]: df_reduced_encoded.head()

```

```

Out[41]:
   total_items  subtotal  num_distinct_items  min_item_price  max_item_price  total_onshift_
0           4      3441             4           557           1239
1           1      1900             1           1400           1400
2           1      1900             1           1900           1900
3           6      6900             5            600           1800
4           3      3900             3           1100           1600

```

```

In [42]: # Split data into features and target
y = df_reduced_encoded['delivery_time_mins']
X = df_reduced_encoded.drop('delivery_time_mins', axis=1)

# Split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale features
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

```

In [132]: from tensorflow.keras.callbacks import LearningRateScheduler
from keras.layers import LeakyReLU
from tensorflow.keras.losses import LogCosh, Huber
from tensorflow.keras.callbacks import Callback
import tensorflow as tf
from tensorflow.keras.losses import MeanSquaredError

```

```

In [159]: # Model definition function
def create_and_train_model(X_train, y_train, X_test, y_test, epochs=50):

    # Model definition
    def create_model(input_shape):
        model = Sequential([
            Dense(32, input_shape=(input_shape,)), kernel_initializer='he_normal'),
            BatchNormalization(),

```

```

        LeakyReLU(alpha=0.01),
        Dense(64),
        BatchNormalization(),
        LeakyReLU(alpha=0.01),
        Dense(128),
        BatchNormalization(),
        LeakyReLU(alpha=0.01),
        Dense(64),
        BatchNormalization(),
        LeakyReLU(alpha=0.01),
        Dense(1, activation='linear')
    ])

    model.compile(optimizer=Adam(learning_rate=0.001, global_clipnorm=1.0), loss='mse')
    return model

class GradientLogger(Callback):
    def __init__(self, train_data):
        super().__init__()
        self.train_data = train_data # Training data as (X, y)

    def on_epoch_end(self, epoch, logs=None):
        # Get the training data
        X, y = self.train_data
        with tf.GradientTape() as tape:
            # Forward pass
            y_pred = self.model(X, training=True) # Access self.model directly
            # Compute the loss
            loss = self.model.compiled_loss(y, y_pred)

        # Compute gradients
        gradients = tape.gradient(loss, self.model.trainable_variables)

        # Log gradient statistics
        grad_norms = [tf.norm(g).numpy() for g in gradients if g is not None]
        print(f"\nEpoch {epoch + 1}: Gradient Norms:")
        print(f"    Mean: {np.mean(grad_norms):.4f}")
        print(f"    Min: {np.min(grad_norms):.4f}")
        print(f"    Max: {np.max(grad_norms):.4f}")

X_train_sample, y_train_sample = X_train[:100], y_train[:100]
gradient_logger = GradientLogger(train_data=(X_train_sample, y_train_sample))

# Create and train the model
input_shape = X_train.shape[1]
model = create_model(input_shape)
history = model.fit(X_train, y_train, epochs=epochs, verbose=1, validation_split=0.1)

# Predictions and final MSE calculation
predictions = model.predict(X_test)
mse = mean_squared_error(y_test, predictions)

# Plot training and validation loss
plt.figure(figsize=(10, 4))
plt.plot(history.history['loss'], label='Training Loss', color='blue')
plt.plot(history.history['val_loss'], label='Validation Loss', color='orange')

```

```

plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.title('Training and Validation Loss Across Epochs')
plt.legend()
plt.grid(True)
plt.show()

return mse

```

```

In [160... # Train model with all features
mse_all_features = create_and_train_model(X_train_scaled, y_train, X_test_scaled, y

```

Epoch 1/50

C:\ProgramData\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)
C:\ProgramData\anaconda3\Lib\site-packages\keras\src\layers\activations\leaky_relu.py:41: UserWarning: Argument `alpha` is deprecated. Use `negative_slope` instead.
warnings.warn(

205/216 ————— 0s 4ms/step - loss: 2249.5393 - mae: 45.0351

Epoch 1: Gradient Norms:

Mean: 44.9948

Min: 0.0000

Max: 345.6995


216/216 ————— 4s 6ms/step - loss: 2238.4680 - mae: 44.9097 - val_loss: 1803.5947 - val_mae: 39.8345


Epoch 2/50


26/216 ————— 0s 4ms/step - loss: 1605.2399 - mae: 37.4779


C:\ProgramData\anaconda3\Lib\site-packages\keras\src\backend\tensorflow\trainer.py:64: UserWarning: `model.compiled_loss()` is deprecated. Instead, use `model.compute_loss(x, y, y_pred, sample_weight, training)`.


warnings.warn(


203/216  **0s** 4ms/step - loss: 1440.5004 - mae: 35.1261
Epoch 2: Gradient Norms:
Mean: 56.2747
Min: 0.0000
Max: 322.0479


216/216  **1s** 5ms/step - loss: 1426.9658 - mae: 34.9186 - val_loss: 871.9685 - val_mae: 26.0319
Epoch 3/50


215/216  **0s** 4ms/step - loss: 613.4295 - mae: 20.5820
Epoch 3: Gradient Norms:
Mean: 54.2441
Min: 0.0000
Max: 196.7765


216/216  **1s** 5ms/step - loss: 612.0215 - mae: 20.5487 - val_loss: 250.4935 - val_mae: 11.8398
Epoch 4/50


212/216  **0s** 4ms/step - loss: 200.9866 - mae: 10.7713
Epoch 4: Gradient Norms:
Mean: 49.0397
Min: 0.0000
Max: 246.7629


216/216  **1s** 5ms/step - loss: 200.6886 - mae: 10.7671 - val_loss: 190.4771 - val_mae: 10.5099
Epoch 5/50


214/216  **0s** 4ms/step - loss: 178.1042 - mae: 10.4803
Epoch 5: Gradient Norms:
Mean: 43.3286
Min: 0.0000
Max: 220.3572


216/216  **1s** 5ms/step - loss: 178.0918 - mae: 10.4800 - val_loss: 183.3919 - val_mae: 10.3654
Epoch 6/50


213/216  **0s** 4ms/step - loss: 174.6668 - mae: 10.3751
Epoch 6: Gradient Norms:
Mean: 42.3577
Min: 0.0000
Max: 232.2325


216/216  **1s** 5ms/step - loss: 174.6826 - mae: 10.3756 - val_loss: 181.3274 - val_mae: 10.3360
Epoch 7/50


216/216  **0s** 5ms/step - loss: 174.4046 - mae: 10.3652
Epoch 7: Gradient Norms:
Mean: 41.7513
Min: 0.0000
Max: 227.0512


216/216  **1s** 5ms/step - loss: 174.4054 - mae: 10.3653 - val_loss: 179.9899 - val_mae: 10.3920
Epoch 8/50


211/216  **0s** 4ms/step - loss: 173.5091 - mae: 10.3272
Epoch 8: Gradient Norms:
Mean: 39.6308
Min: 0.0000
Max: 195.9327


216/216  **1s** 5ms/step - loss: 173.5175 - mae: 10.3279 - val_loss: 178.9618 - val_mae: 10.3840
Epoch 9/50


213/216  **0s** 4ms/step - loss: 172.7879 - mae: 10.3252
Epoch 9: Gradient Norms:
Mean: 37.7135
Min: 0.0000
Max: 194.6149


216/216  **1s** 5ms/step - loss: 172.7924 - mae: 10.3252 - val_loss: 177.3396 - val_mae: 10.3103
Epoch 10/50


206/216  **0s** 4ms/step - loss: 172.7053 - mae: 10.3070
Epoch 10: Gradient Norms:
Mean: 36.7855
Min: 0.0000
Max: 201.8718


216/216  **1s** 5ms/step - loss: 172.6942 - mae: 10.3072 - val_loss: 176.8672 - val_mae: 10.3415
Epoch 11/50


212/216  **0s** 4ms/step - loss: 169.8870 - mae: 10.1976
Epoch 11: Gradient Norms:
Mean: 38.6563
Min: 0.0000
Max: 217.1819


216/216  **1s** 5ms/step - loss: 169.9298 - mae: 10.1994 - val_loss: 175.9690 - val_mae: 10.3145
Epoch 12/50


206/216  **0s** 4ms/step - loss: 170.0631 - mae: 10.2443
Epoch 12: Gradient Norms:
Mean: 38.4658
Min: 0.0000
Max: 211.4359


216/216  **1s** 5ms/step - loss: 170.1384 - mae: 10.2463 - val_loss: 175.5651 - val_mae: 10.3350
Epoch 13/50


212/216  **0s** 4ms/step - loss: 171.0560 - mae: 10.2680
Epoch 13: Gradient Norms:
Mean: 36.0993
Min: 0.0000
Max: 200.1110


216/216  **1s** 5ms/step - loss: 171.0535 - mae: 10.2677 - val_loss: 176.2726 - val_mae: 10.4024
Epoch 14/50


206/216  **0s** 4ms/step - loss: 169.7702 - mae: 10.2290
Epoch 14: Gradient Norms:
Mean: 38.5769
Min: 0.0000
Max: 215.0470


216/216  **1s** 5ms/step - loss: 169.8285 - mae: 10.2305 - val_loss: 176.0461 - val_mae: 10.2768
Epoch 15/50


212/216  **0s** 4ms/step - loss: 169.8562 - mae: 10.2204
Epoch 15: Gradient Norms:
Mean: 37.8294
Min: 0.0000
Max: 226.3993


216/216  **1s** 5ms/step - loss: 169.8674 - mae: 10.2209 - val_loss: 175.1677 - val_mae: 10.3324
Epoch 16/50


209/216  **0s** 4ms/step - loss: 169.8015 - mae: 10.2303
Epoch 16: Gradient Norms:
Mean: 35.2032
Min: 0.0000
Max: 197.5538


216/216  **1s** 5ms/step - loss: 169.8246 - mae: 10.2307 - val_loss: 175.3601 - val_mae: 10.3170
Epoch 17/50


210/216  **0s** 5ms/step - loss: 168.2390 - mae: 10.1704
Epoch 17: Gradient Norms:
Mean: 35.2649
Min: 0.0000
Max: 199.6725


216/216  **1s** 5ms/step - loss: 168.2965 - mae: 10.1725 - val_loss: 173.9087 - val_mae: 10.3053
Epoch 18/50


211/216  **0s** 5ms/step - loss: 169.8642 - mae: 10.2347
Epoch 18: Gradient Norms:
Mean: 37.6104
Min: 0.0000
Max: 215.0655


216/216  **1s** 5ms/step - loss: 169.8629 - mae: 10.2344 - val_loss: 176.5241 - val_mae: 10.4087
Epoch 19/50


207/216  **0s** 4ms/step - loss: 170.5861 - mae: 10.2477
Epoch 19: Gradient Norms:
Mean: 36.8726
Min: 0.0000
Max: 219.9571


216/216  **1s** 5ms/step - loss: 170.5537 - mae: 10.2464 - val_loss: 173.9826 - val_mae: 10.2936
Epoch 20/50


204/216  **0s** 4ms/step - loss: 169.0736 - mae: 10.2095
Epoch 20: Gradient Norms:
Mean: 36.7550
Min: 0.0000
Max: 212.5412


216/216  **1s** 5ms/step - loss: 169.0885 - mae: 10.2091 - val_loss: 174.1143 - val_mae: 10.3169
Epoch 21/50


201/216  **0s** 4ms/step - loss: 169.0007 - mae: 10.2054
Epoch 21: Gradient Norms:
Mean: 42.3597
Min: 0.0000
Max: 248.8876


216/216  **1s** 5ms/step - loss: 169.0352 - mae: 10.2057 - val_loss: 175.1497 - val_mae: 10.4070
Epoch 22/50


209/216  **0s** 5ms/step - loss: 170.0843 - mae: 10.2112
Epoch 22: Gradient Norms:
Mean: 37.9089
Min: 0.0000
Max: 232.0540


216/216  **1s** 5ms/step - loss: 170.0564 - mae: 10.2109 - val_loss: 174.8790 - val_mae: 10.3625
Epoch 23/50


206/216  **0s** 4ms/step - loss: 168.7414 - mae: 10.1781
Epoch 23: Gradient Norms:
Mean: 37.8884
Min: 0.0000
Max: 223.5155


216/216  **1s** 5ms/step - loss: 168.7738 - mae: 10.1795 - val_loss: 174.5758 - val_mae: 10.3420
Epoch 24/50


213/216  **0s** 4ms/step - loss: 170.0041 - mae: 10.2206
Epoch 24: Gradient Norms:
Mean: 35.8177
Min: 0.0000
Max: 204.6021


216/216  **1s** 5ms/step - loss: 169.9845 - mae: 10.2201 - val_loss: 173.8655 - val_mae: 10.3372
Epoch 25/50


209/216  **0s** 4ms/step - loss: 167.3510 - mae: 10.1334
Epoch 25: Gradient Norms:
Mean: 36.8118
Min: 0.0000
Max: 213.9263


216/216  **1s** 5ms/step - loss: 167.4093 - mae: 10.1354 - val_loss: 173.1388 - val_mae: 10.2854
Epoch 26/50


214/216  **0s** 4ms/step - loss: 167.1989 - mae: 10.1519
Epoch 26: Gradient Norms:
Mean: 38.1846
Min: 0.0000
Max: 226.2578


216/216  **1s** 5ms/step - loss: 167.2194 - mae: 10.1523 - val_loss: 174.7395 - val_mae: 10.3924
Epoch 27/50


213/216  **0s** 4ms/step - loss: 168.4286 - mae: 10.1805
Epoch 27: Gradient Norms:
Mean: 39.3817
Min: 0.0000
Max: 229.7686


216/216  **1s** 5ms/step - loss: 168.4355 - mae: 10.1807 - val_loss: 173.5632 - val_mae: 10.2985
Epoch 28/50


213/216  **0s** 5ms/step - loss: 169.3706 - mae: 10.1954
Epoch 28: Gradient Norms:
Mean: 34.7444
Min: 0.0000
Max: 193.5687


216/216  **1s** 5ms/step - loss: 169.3543 - mae: 10.1949 - val_loss: 174.1017 - val_mae: 10.2638
Epoch 29/50


215/216  **0s** 4ms/step - loss: 168.0120 - mae: 10.1593
Epoch 29: Gradient Norms:
Mean: 35.0953
Min: 0.0000
Max: 203.3445


216/216  **1s** 5ms/step - loss: 168.0181 - mae: 10.1596 - val_loss: 173.0428 - val_mae: 10.3166
Epoch 30/50


206/216  **0s** 4ms/step - loss: 168.3662 - mae: 10.1597
Epoch 30: Gradient Norms:
Mean: 32.5813
Min: 0.0000
Max: 188.6980


216/216  **1s** 5ms/step - loss: 168.3691 - mae: 10.1603 - val_loss: 173.2359 - val_mae: 10.2668
Epoch 31/50


211/216  **0s** 4ms/step - loss: 168.0315 - mae: 10.1484
Epoch 31: Gradient Norms:
Mean: 36.5598
Min: 0.0000
Max: 203.8364


216/216  **1s** 5ms/step - loss: 168.0353 - mae: 10.1487 - val_loss: 172.8335 - val_mae: 10.2614
Epoch 32/50


208/216  **0s** 4ms/step - loss: 167.6252 - mae: 10.1483
Epoch 32: Gradient Norms:
Mean: 37.2033
Min: 0.0000
Max: 215.8843


216/216  **1s** 5ms/step - loss: 167.6480 - mae: 10.1491 - val_loss: 174.2439 - val_mae: 10.2597
Epoch 33/50


207/216  **0s** 4ms/step - loss: 167.5430 - mae: 10.1612
Epoch 33: Gradient Norms:
Mean: 36.2382
Min: 0.0000
Max: 200.4900


216/216  **1s** 5ms/step - loss: 167.5444 - mae: 10.1605 - val_loss: 174.8259 - val_mae: 10.3340
Epoch 34/50


211/216  **0s** 4ms/step - loss: 167.3783 - mae: 10.1375
Epoch 34: Gradient Norms:
Mean: 35.8597
Min: 0.0000
Max: 198.7228


216/216  **1s** 5ms/step - loss: 167.3972 - mae: 10.1382 - val_loss: 173.4129 - val_mae: 10.3532
Epoch 35/50


208/216  **0s** 5ms/step - loss: 167.4493 - mae: 10.1517
Epoch 35: Gradient Norms:
Mean: 33.2744
Min: 0.0000
Max: 198.6350


216/216  **1s** 5ms/step - loss: 167.4669 - mae: 10.1519 - val_loss: 172.8176 - val_mae: 10.3464
Epoch 36/50


204/216  **0s** 4ms/step - loss: 166.0669 - mae: 10.1092
Epoch 36: Gradient Norms:
Mean: 35.4396
Min: 0.0000
Max: 210.2507


216/216  **1s** 5ms/step - loss: 166.1475 - mae: 10.1111 - val_loss: 173.1565 - val_mae: 10.3404
Epoch 37/50


214/216  **0s** 4ms/step - loss: 167.8537 - mae: 10.1367
Epoch 37: Gradient Norms:
Mean: 35.3704
Min: 0.0000
Max: 202.6881


216/216  **1s** 5ms/step - loss: 167.8489 - mae: 10.1368 - val_loss: 174.0972 - val_mae: 10.1805
Epoch 38/50


207/216  **0s** 4ms/step - loss: 165.7655 - mae: 10.0874
Epoch 38: Gradient Norms:
Mean: 34.7899
Min: 0.0000
Max: 203.2329


216/216  **1s** 5ms/step - loss: 165.8356 - mae: 10.0894 - val_loss: 172.6210 - val_mae: 10.2326
Epoch 39/50


216/216  **0s** 4ms/step - loss: 166.5712 - mae: 10.1160
Epoch 39: Gradient Norms:
Mean: 33.4628
Min: 0.0000
Max: 193.4768


216/216  **1s** 5ms/step - loss: 166.5741 - mae: 10.1161 - val_loss: 174.6069 - val_mae: 10.2312
Epoch 40/50


215/216  **0s** 5ms/step - loss: 167.2635 - mae: 10.1377
Epoch 40: Gradient Norms:
Mean: 36.1214
Min: 0.0000
Max: 199.0746


216/216  **1s** 6ms/step - loss: 167.2632 - mae: 10.1377 - val_loss: 173.2289 - val_mae: 10.2637
Epoch 41/50


214/216  **0s** 4ms/step - loss: 165.9865 - mae: 10.1016
Epoch 41: Gradient Norms:
Mean: 36.1178
Min: 0.0000
Max: 187.8597


216/216  **1s** 5ms/step - loss: 165.9996 - mae: 10.1020 - val_loss: 173.3917 - val_mae: 10.2460
Epoch 42/50


208/216  **0s** 4ms/step - loss: 167.4256 - mae: 10.1377
Epoch 42: Gradient Norms:
Mean: 34.2584
Min: 0.0000
Max: 183.8495


216/216  **1s** 5ms/step - loss: 167.4053 - mae: 10.1371 - val_loss: 172.6636 - val_mae: 10.2706
Epoch 43/50


208/216  **0s** 4ms/step - loss: 166.2100 - mae: 10.0899
Epoch 43: Gradient Norms:
Mean: 38.2909
Min: 0.0000
Max: 221.6710


216/216  **1s** 5ms/step - loss: 166.2339 - mae: 10.0910 - val_loss: 173.8541 - val_mae: 10.3175
Epoch 44/50


215/216  **0s** 4ms/step - loss: 166.2875 - mae: 10.0918
Epoch 44: Gradient Norms:
Mean: 38.1552
Min: 0.0000
Max: 209.5925


216/216  **1s** 5ms/step - loss: 166.2926 - mae: 10.0920 - val_loss: 174.0124 - val_mae: 10.3583
Epoch 45/50


216/216  **0s** 4ms/step - loss: 167.4233 - mae: 10.1576
Epoch 45: Gradient Norms:
Mean: 35.2712
Min: 0.0000
Max: 206.3295


216/216  **1s** 5ms/step - loss: 167.4204 - mae: 10.1574 - val_loss: 173.3062 - val_mae: 10.3218
Epoch 46/50


207/216  **0s** 4ms/step - loss: 164.5789 - mae: 10.0513
Epoch 46: Gradient Norms:
Mean: 31.8549
Min: 0.0000
Max: 164.0808


216/216  **1s** 5ms/step - loss: 164.6734 - mae: 10.0542 - val_loss: 173.2081 - val_mae: 10.3170
Epoch 47/50


209/216  **0s** 4ms/step - loss: 166.5354 - mae: 10.1125
Epoch 47: Gradient Norms:
Mean: 35.1648
Min: 0.0000
Max: 205.3537


216/216  **1s** 5ms/step - loss: 166.5397 - mae: 10.1125 - val_loss: 173.4931 - val_mae: 10.3266
Epoch 48/50


206/216  **0s** 4ms/step - loss: 166.5779 - mae: 10.1161
Epoch 48: Gradient Norms:
Mean: 32.0254
Min: 0.0000
Max: 173.7895


216/216  **1s** 5ms/step - loss: 166.5529 - mae: 10.1153 - val_loss: 173.8577 - val_mae: 10.3471
Epoch 49/50


212/216  **0s** 4ms/step - loss: 164.1559 - mae: 10.0640
Epoch 49: Gradient Norms:
Mean: 36.6529
Min: 0.0000
Max: 195.7887

216/216  **1s** 5ms/step - loss: 164.2043 - mae: 10.0651 - val_loss: 173.3613 - val_mae: 10.2497
Epoch 50/50

211/216  **0s** 4ms/step - loss: 166.1796 - mae: 10.0880
Epoch 50: Gradient Norms:
Mean: 32.6656
Min: 0.0000
Max: 177.2888

216/216  **1s** 5ms/step - loss: 166.1765 - mae: 10.0882 - val_loss: 172.7631 - val_mae: 10.2411

960/960  **1s** 983us/step



In [161... mse_all_features

Out[161... 168.559646042214

- Increasing the number of layers, number of neurons or changing the optimizers did not give better performance than the above model